

Lecture 3: Logistic regression


Classification

- Email: Spam / Not Spam?
- Online Transactions: Fraudulent (Yes / No)?
- Tumor: Malignant / Benign ?


$$y \in \{0, 1\}$$

0: “Negative Class” (e.g., benign tumor)

1: “Positive Class” (e.g., malignant tumor)



Threshold classifier output



Perceptron

A perceptron, first introduced in 1958, is a deterministic binary classifier of the following form:


$$h_{\theta}(x) = \mathbb{I}(w^T x + b > 0)$$

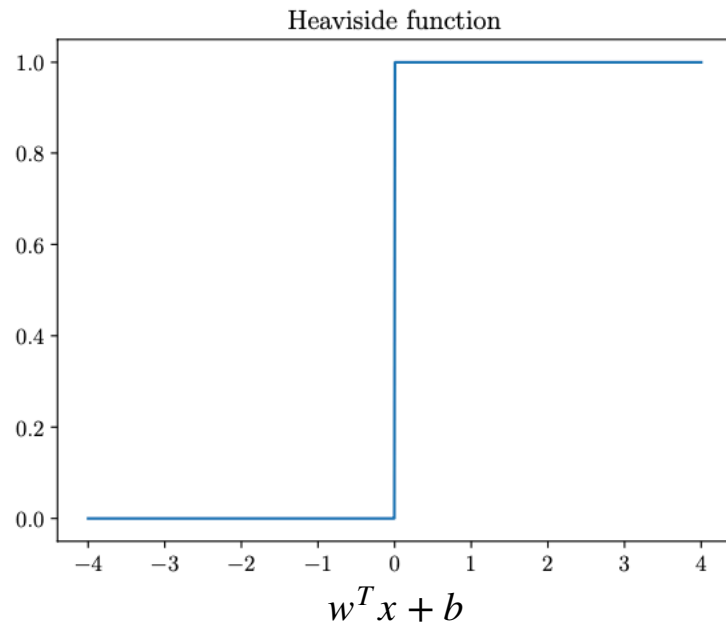
Recall that $w^T x = [w_0 \quad w_1 \quad w_3 \quad \cdots] \begin{bmatrix} x_0 \\ x_1 \\ x_3 \\ \vdots \end{bmatrix} = \sum_i w_i x_i$

b : bias term

\mathbb{I} : Heaviside function

Decision boundary: a linear 'subspace' where $w^T x + b = 0$





$$h_{\theta}(x) = \mathbb{I}(w^T x + b > 0)$$

Probabilistic formulation

- Bernoulli distribution for binary random variables (e.g. outcomes of flipping a coin)
- ψ : probability of getting the outcome '1' (malignant)

$$\text{Ber}(y | \psi) = \begin{cases} 1 - \psi & \text{if } y = 0 \\ \psi & \text{if } y = 1 \end{cases}$$

- This can be concisely expressed as

$$\text{Ber}(y | \psi) = \psi^y (1 - \psi)^{1-y}$$

- Model the probability ψ as a function of the input X

- Perceptron: $y = 0$ or 1

$h_{\theta}(x)$ can be 1 or 0



- Logistic Regression: $0 \leq h_{\theta}(x) \leq 1$

Logistic function

- We want $0 \leq h_{\theta}(x) \leq 1$
- We do this using the logistic function

$$h_{\theta}(x) = g(w^T x + b) = g(\theta^T x)$$

$$g(z) = \frac{1}{1 + e^{-z}}$$



Interpretation of Hypothesis Output


$h_{\theta}(x)$ = estimated probability that $y = 1$ on input x

Example: If $x = \begin{bmatrix} x_0 \\ x_1 \end{bmatrix} = \begin{bmatrix} 1 \\ \text{tumorSize} \end{bmatrix}$

$$h_{\theta}(x) = 0.7$$


Tell patient that 70% chance of tumor being malignant

Decision Boundary



$$h_{\theta}(x) = g(\theta^T x) \quad h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$$

$$g(z) = \frac{1}{1+e^{-z}}$$



predict “ $y = 1$ ” if $h_{\theta}(x) \geq 0.5$

predict “ $y = 0$ ” if $h_{\theta}(x) < 0.5$

Training set: $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$

m examples $x \in \begin{bmatrix} x_0 \\ x_1 \\ \dots \\ x_n \end{bmatrix} \quad x_0 = 1, y \in \{0, 1\}$


$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

How to choose the parameters θ ?

Cost function

$$\text{Linear regression: } J(\theta) = \frac{1}{m} \sum_{i=1}^m \frac{1}{2} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

$$\text{Cost}(h_{\theta}(x^{(i)}), y^{(i)}) = \frac{1}{2} (h_{\theta}(x^{(i)}) - y^{(i)})^2$$



Logistic regression cost function - Likelihood

Likelihood: probability of observing the data given the model parameters

Likelihood for observing data point i :

$$\text{Ber}(y_i | h_{\theta}(x_i)) = h_{\theta}(x_i)^{y_i} \times (1 - h_{\theta}(x_i))^{1-y_i}$$

$$\text{Ber}(y | \psi) = \begin{cases} 1 - \psi & \text{if } y = 0 \\ \psi & \text{if } y = 1 \end{cases}$$

$$\text{Ber}(y | \psi) = \psi^y (1 - \psi)^{1-y}$$

Logistic regression cost function - Likelihood

Likelihood for observing data point i :

$$\text{Ber}(y_i | h_{\theta}(x_i)) = h_{\theta}(x_i)^{y_i} \times (1 - h_{\theta}(x_i))^{1-y_i}$$

Likelihood for observing all m data points

$$\prod_i^m h_{\theta}(x_i)^{y_i} \times (1 - h_{\theta}(x_i))^{1-y_i}$$

Logistic regression cost function - Likelihood

Likelihood for observing all m data points

$$\prod_i^m h_{\theta}(x_i)^{y_i} \times (1 - h_{\theta}(x_i))^{1-y_i}$$

We usually work with the negative log likelihood

$$\text{NLL}(\theta) = -\frac{1}{m} \sum_i^m [y_i \log h_{\theta}(x_i) + (1 - y_i) \log(1 - h_{\theta}(x_i))]$$

Maximum likelihood

We find the best model parameters $\hat{\theta}$ by maximizing the NLL

$$\text{NLL}(\theta) = -\frac{1}{m} \sum_i^m [y_i \log h_{\theta}(x_i) + (1 - y_i) \log(1 - h_{\theta}(x_i))]$$

$$\hat{\theta} = \operatorname{argmin}_{\theta} -\frac{1}{m} \sum_i^m [y_i \log h_{\theta}(x_i) + (1 - y_i) \log(1 - h_{\theta}(x_i))]$$

No close form solution exists for the logistic regression cost

$$\begin{aligned} J(\theta) &= \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)}) \\ &= -\frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(x^{(i)})) \right] \end{aligned}$$

To fit parameters θ :

$$\min_{\theta} J(\theta)$$

To make a prediction given new x :

$$\text{Output } h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

Gradient Descent

$$J(\theta) = -\frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(x^{(i)})) \right]$$

Want $\min_{\theta} J(\theta)$:

Repeat {

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

}

(simultaneously update all θ_j)

Gradient Descent

$$J(\theta) = -\frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(x^{(i)})) \right]$$

Want $\min_{\theta} J(\theta)$:

Repeat {

$$\theta_j := \theta_j - \alpha \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$


} (simultaneously update all θ_j)

Algorithm looks identical to linear regression!

Negative log likelihood is equal to **binary cross entropy**

- The entropy of a distribution

$$\mathbb{H}(p) = - \sum_x p(x) \log(p(x))$$




Negative log likelihood is equal to **binary cross entropy**

- The cross entropy between two probability distributions

$$\mathbb{H}(p, q) = - \sum_x p(x) \log(q(x))$$

- Cross entropy for two Bernoulli distributions

$$\mathbb{H}(p, q) = - \left[\psi_p \log \psi_q + (1 - \psi_p) \log(1 - \psi_q) \right]$$



cross entropy of $\text{Ber}(\psi)$ relative to $\text{Ber}(0.3)$

Negative log likelihood is equal to **binary cross entropy**

- Cross entropy for two Bernoulli distributions
- Consider the data point i as the distribution p
- Distribution q is the modeled distribution with probability $h_{\theta}(x_i)$

$$\mathbb{H}(p, q) = - \left[\psi_p \log \psi_q + (1 - \psi_p) \log(1 - \psi_q) \right]$$

$$\mathbb{H}(p, q) = - \left[y_i \log h_{\theta}(x_i) + (1 - y_i) \log(1 - h_{\theta}(x_i)) \right]$$

$$\text{NLL}(\theta) = - \frac{1}{m} \sum_i^m \mathbb{H}(y_i, h_{\theta}(x_i))$$

$$\text{NLL}(\theta) = - \frac{1}{m} \sum_i^m \left[y_i \log h_{\theta}(x_i) + (1 - y_i) \log(1 - h_{\theta}(x_i)) \right]$$

Negative log likelihood is equal to **binary cross entropy**

- Cross entropy also equal to the entropy of p and the KL divergence between p and q

$$\mathbb{H}(p, q) = H(p) + D_{KL}(p \| q)$$

$$\begin{aligned}\mathbb{H}(p, q) &= - \sum_x p(x) \log(q(x)) \\ &= - \sum_x p(x) (\log(q(x)) + \log(p(x)) - \log(p(x))) \\ &= - \sum_x p(x) \log(p(x)) - \sum_x p(x) (\log(q(x)) - \log(p(x))) \\ &= \mathbb{H}(p) + \sum_x p(x) \frac{\log(p(x))}{\log(q(x))}\end{aligned}$$

$$D_{KL}(p \| q) = \sum_x p(x) \log \left(\frac{p(x)}{q(x)} \right)$$

Multiclass case


Multiclass classification

Email foldering/tagging: Work, Friends, Family, Hobby


Medical diagrams: Not ill, Cold, Flu

Weather: Sunny, Cloudy, Rain, Snow

Binary classification:



Multi-class classification:



- Recall in logistic regression

$$h_{\theta}(x) = g(\theta^T x)$$

$$g(z) = \frac{1}{1+e^{-z}}$$

This can be considered as a procedure where we assign two numbers to the two outcomes

$$y = 0 \leftarrow 0$$

$$y = 1 \leftarrow \theta^T x$$

And renormalized such that their sum is one

$$p(y = 0) = \frac{e^0}{e^0 + e^{\theta^T x}} \text{ and } p(y = 1) = \frac{e^{\theta^T x}}{e^0 + e^{\theta^T x}} = \frac{1}{e^{-\theta^T x} + 1}$$

- The **softmax** function generalizes the logistic function

$$\sigma(\mathbf{z})_i = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}} .$$

- We can build a model for multi-label classification using the softmax function


$$a = \begin{bmatrix} a_0 \\ a_1 \\ \vdots \\ a_{C-1} \end{bmatrix} = \begin{bmatrix} 0 \\ w_1^T x \\ \vdots \\ w_{C-1}^T x \end{bmatrix} = Wx$$

W is the $C \times D$ dimensional weight matrix
 x is the D dimensional input vector

- And model the output probability using


$$p(y = c | \mathbf{x}, \boldsymbol{\theta}) = \frac{e^{a_c}}{\sum_{c'=1}^C e^{a_{c'}}}$$

Non-linear decision boundaries



$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2)$$

Predict “ $y = 1$ ” if $-1 + x_1^2 + x_2^2 \geq 0$



$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_1^2 x_2 + \theta_5 x_1^2 x_2^2 + \theta_6 x_1^3 x_2 + \dots)$$